


Docker

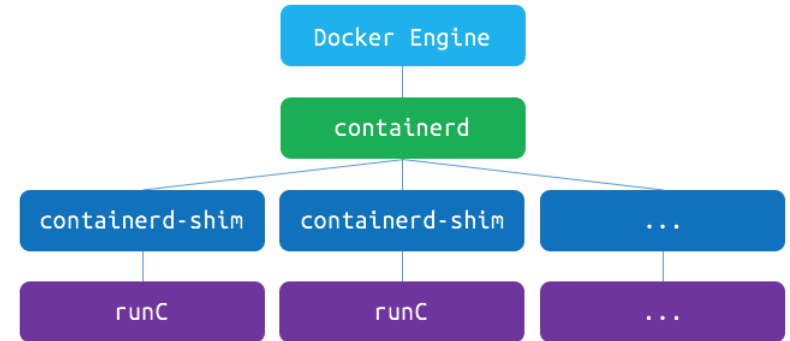
Docker

- Broadly speaking, Docker is a platform for developing, shipping and running applications using container technology
 - Founded in 2009.
 - Formerly dotCloud Inc.
 - Released Docker in 2013.
- It consists of a bunch of products/tools
 - Docker Engine – i.e., to start container instances
 - Docker Hub – i.e., like github to host public container images
 - Docker Trusted Registry – i.e., store container images
 - Docker Machine - i.e., create a (virtual) machine that support docker container
 - Docker Compose – i.e., to build container images
 - Docker for Windows/Mac
 - Docker Datacenter (swarm) – i.e., similar to Kubernetes



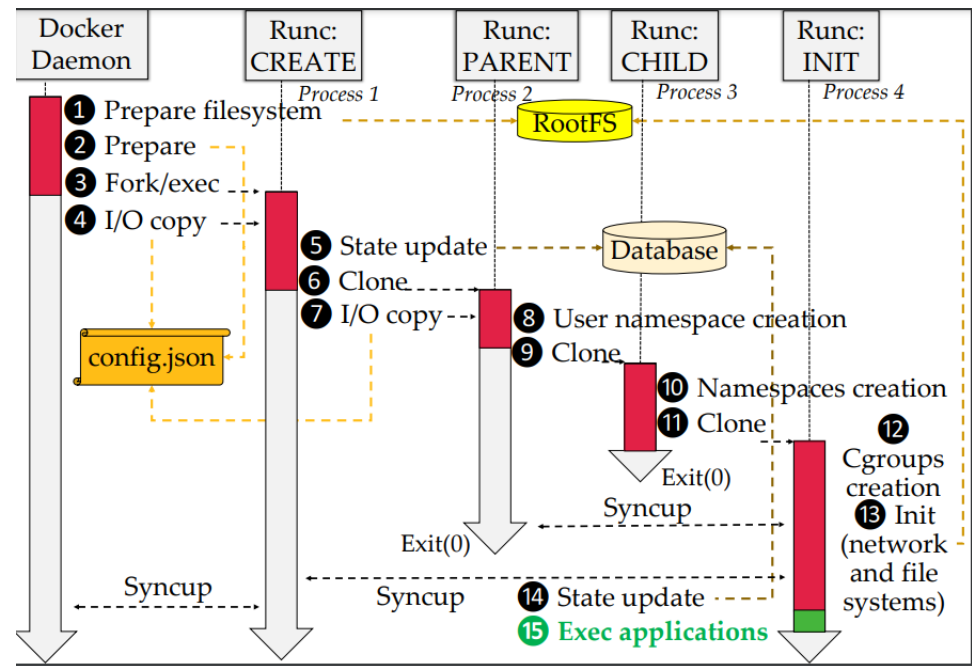
Docker architecture

- Docker Engine receives requests from upstream clients
- Containerd manages the complete container lifecycle
 - Create, pause, termination, deletion
- runC is a lightweight tool that does one thing, it creates a container instance (name spaces and cgroups): <https://github.com/opencontainers/runc>



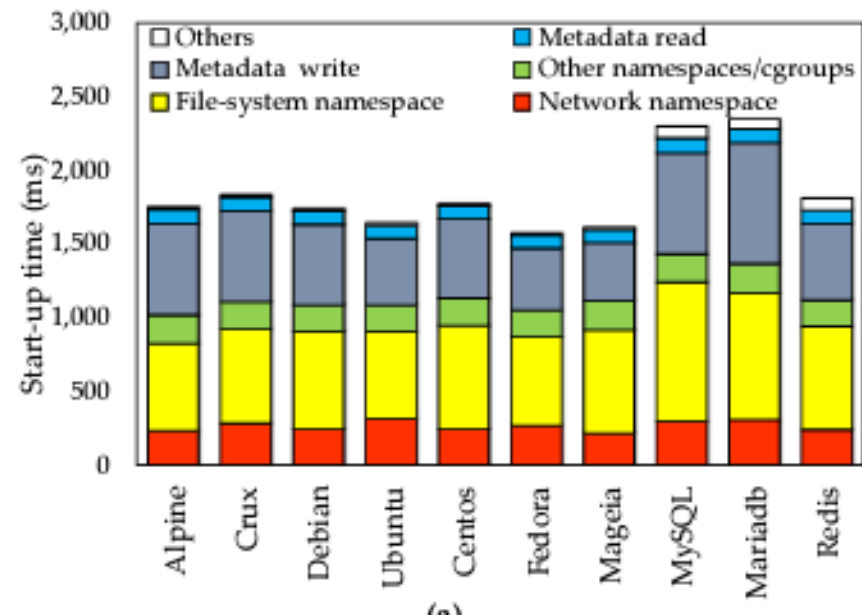
Long startup latency

- The construction of a Docker container goes through a long, serialized pipeline involving multiple processes
- These processes need frequent (and slow) synchronization to coordinate the different initialization stages
- E.g., allocating storage and network resources, isolating allocated resources, and filtering system calls.



Long startup latency

- It could take up to 2.4 seconds to complete a single Linux container initialization before its encapsulated function code executes.
- Creating various isolation components for a container instance contributes more than 50% to the total cold-start latency.




Container Images

- A container instance is launched from a container image.
- A container image is a root file system (mini OS) that includes everything needed to run an application(s)
 - The application code, a runtime, libraries, environment variables, and configuration files.
 - Consisting of folders and files just like a Linux file system (i.e., file organization)
- When we launch a container, a container instance is a runtime instance of an image
 - like binary code vs. processes

Layout of a Docker Container Image


- A container image contains an image base layer, including a base file system.
- When you launch a container instance from the container image, another layer is created on top of the base image layer, called container layer
- Container layer are initially empty and will be discarded when the container instance is terminated
- So all modifications during container execution will be forgotten



- The image base is **read only**
 - Multiple containers can share the same image base layer
 - Just like a shared library
- The container layer can both **read and write**
 - Per-instance, private layer


Write policy

- Writing to a file for the first time and the file exists in the image layer
- `copy_up`: copy files from the base layer to the container layer, and write changes to it.
- Deleting a file
 - A “whiteout” file is created in the container layer marking that the file with the same name in the image layer is invalid



Read Policies

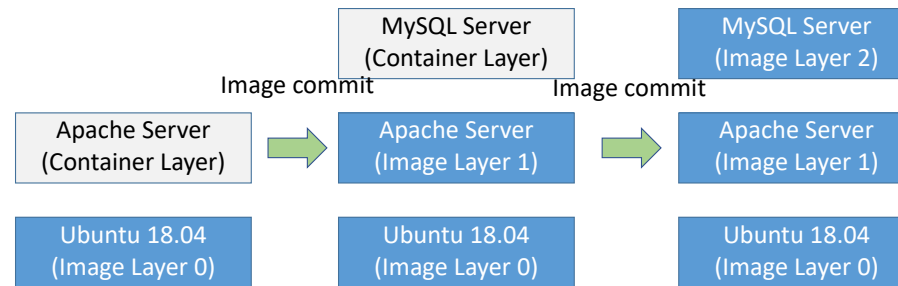
- Files only exist in image layer, it is read from image layer
- Files only exist in container layer, it is read from container layer
- Files exist in both layers, it is read from container
- Files in the container layer obscure files with the same name in the image layer.



Pros/cons of overlay file systems

- Cons
 - Overhead
- Pros
 - Many container instances share the same base images
 - Saving space
 - Container image can be stackable
 - Easy to build new images


Stackable Container images



gVisor


Improving container security by intercepting system calls

Security concern of containers



- Traditional containers have a wide system call interface
 - E.g. Linux has over 400 system calls
- Vulnerable system calls can allow unauthorized access across containers, hosts or data centers etc., thus affecting all the containers on the Host OS.


Option 1



- Allow only limited system calls that the container needs.
- Rule-based execution allows the specification of a fine-grained security policy for an application or container. (e.g., Linux's seccomp)
- In practice, not easy. It may break unknown applications whose system call profile is not known accurately.


gVisor -- Google's Secure Containerization

- gVisor intercepts application system calls and acts as a guest kernel
- It implements a substantial portion of the Linux system surface
- The isolation boundary between the application and the host kernel is maintained
- Drawback: High per-system call overhead



gVisor Architecture


- Sentry: the largest component of gVisor
- Can be thought of as a userspace OS kernel, implementing all the kernel functionality needed by the untrusted application
- System calls are redirected to Sentry
- Sentry will make some host system calls to support its operation, but it will not allow the application to directly control the system calls it makes.



gVisor - Sentry

- gVisor requires a way to implement interception of syscalls
- The `ptrace()` system call provides a mechanism by which a parent process may observe and control the execution of another process.

```
long ptrace(enum __ptrace_request request, pid_t pid, void * addr, void * data);
```



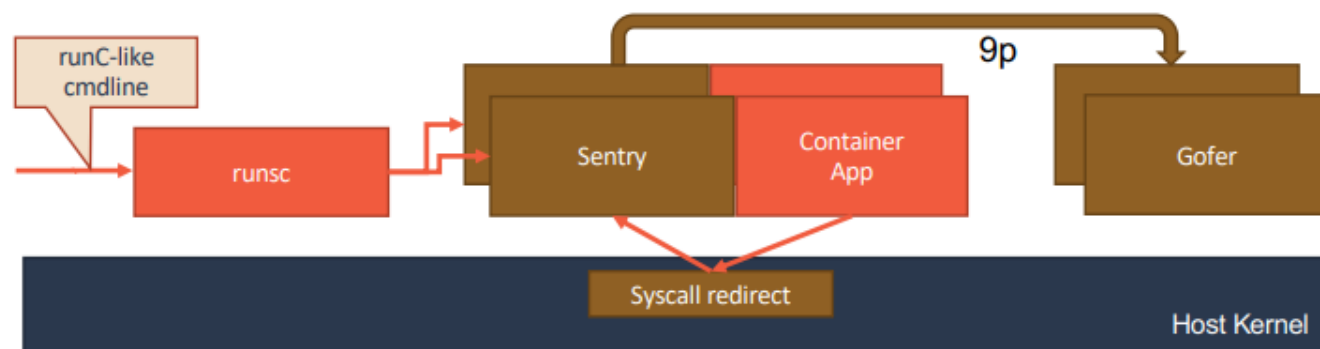
- `PTRACE_SYSEMU` request
 - causes the traced process to stop on entry to the next syscall

gVisor - Sentry

```
for (;;) {  
    ptrace(PTRACE_SYSEMU, pid, 0, 0);  
    waitpid(pid, 0, 0);  
  
    struct user_regs_struct regs;  
    ptrace(PTRACE_GETREGS, pid, 0, &regs);  
  
    switch (regs.orig_rax) {  
        case OS_read:  
            /* ... */  
  
        case OS_write:  
            /* ... */  
  
        case OS_open:  
            /* ... */  
  
        case OS_exit:  
            /* ... */  
  
        /* ... and so on ... */  
    }  
}
```

gVisor - runsc

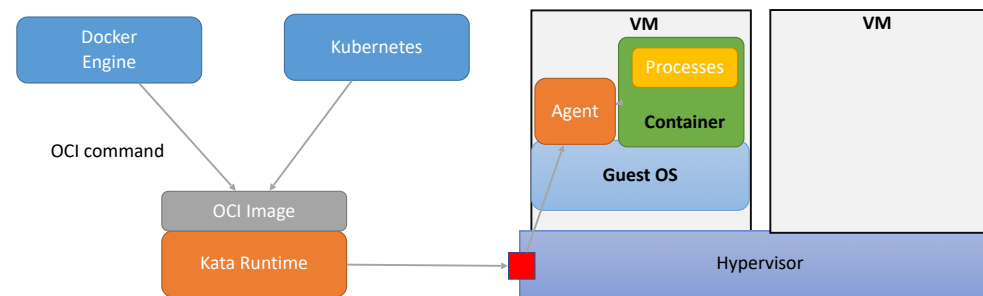
- runsc – the entrypoint to running a sandboxed container
- Implements an OCI runtime specification including:
 - A config.json file contains container configurations
 - A root filesystem



Kata Containers

- Basic Ideas: It's possible to run containers inside of virtual machines
 - Pretty common deployment method (think about the cloud)
 - Introduces another layer of protection: the hypervisor
 - But notice that Hypervisors can have security bugs as well
- Kata Container
 - Introduced in 2017 from the merger of Intel's Clear Containers and Hyper's runV
 - "Wraps" containers into dedicated virtual machines
 - OCI runtime implementation: can be plugged into the container engine (e.g., Docker)
 - Supports existing container images

Kata containers - architecture



References

- 1. How to Implement Secure Containers Using Google's gVisor <https://thenewstack.io/how-to-implement-secure-containers-using-googles-gvisor/>
- 2. Intercepting and Emulating Linux System Calls with Ptrace: <https://nullprogram.com/blog/2018/06/23/>
- 3. The True Cost of Containing: A gVisor Case Study: <https://www.usenix.org/system/files/hotcloud19-paper-young.pdf>
- 4. kata containers & gVisor: <https://object-storage-ca-ymq-1.vexxhost.net/swift/v1/6e4619c416ff4bd19e1c087f27a43eea/www-assets-prod/presentation-media/kata-containers-and-gvisor-a-quantitative-comparison.pdf>
- 5. Kata Containers The way to run virtualized containers: <https://object-storage-ca-ymq-1.vexxhost.net/swift/v1/6e4619c416ff4bd19e1c087f27a43eea/www-assets-prod/presentation-media/Kata-Containers-The-way-to-run-virtualized-containers.pdf>
- 6. Bringing container security to the next level using Kata Containers: https://www.suse.com/media/presentation/TUT1201_Bringing_Container_Security_to_the_Next_Level_Using_Kata_Containers.pdf