

# CS-452/552 Introduction to Cloud Computing

Storage Virtualization

# Data Storage Systems

Data can be stored in various places in different manners

- Hardware: CPU registers, caches, main memory and persistent storage
- Software: File systems, object storage, databases (SQL databases and No-SQL databases).

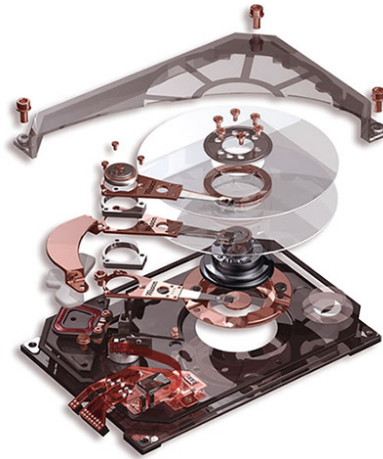


# Storage I/O system within a single host

Persistent Storage media



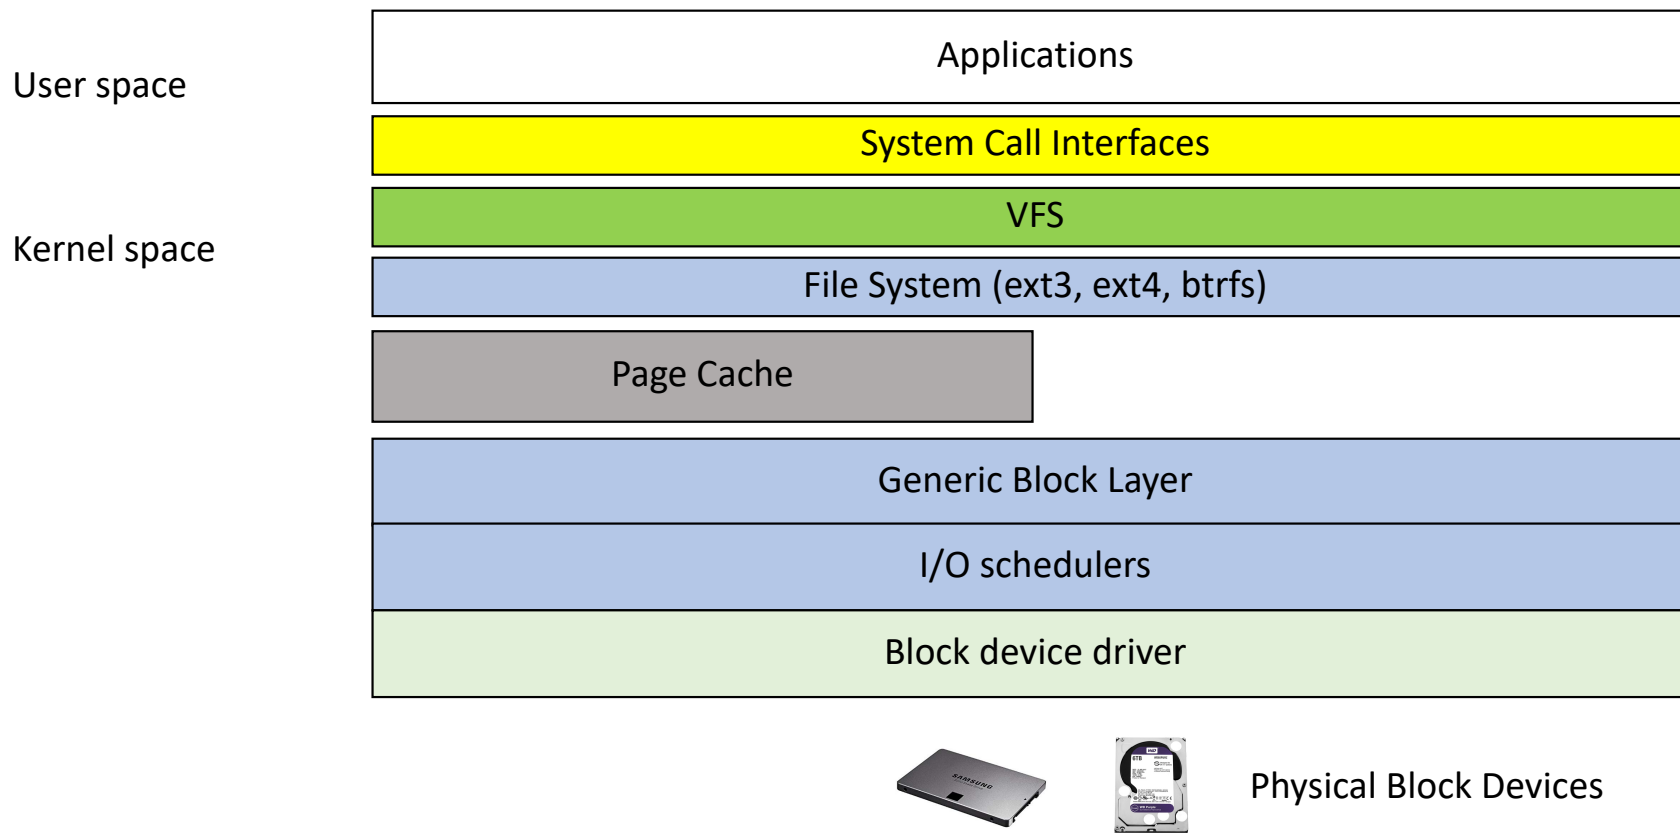
**FLASH**



**HDD OR DISK DRIVE**

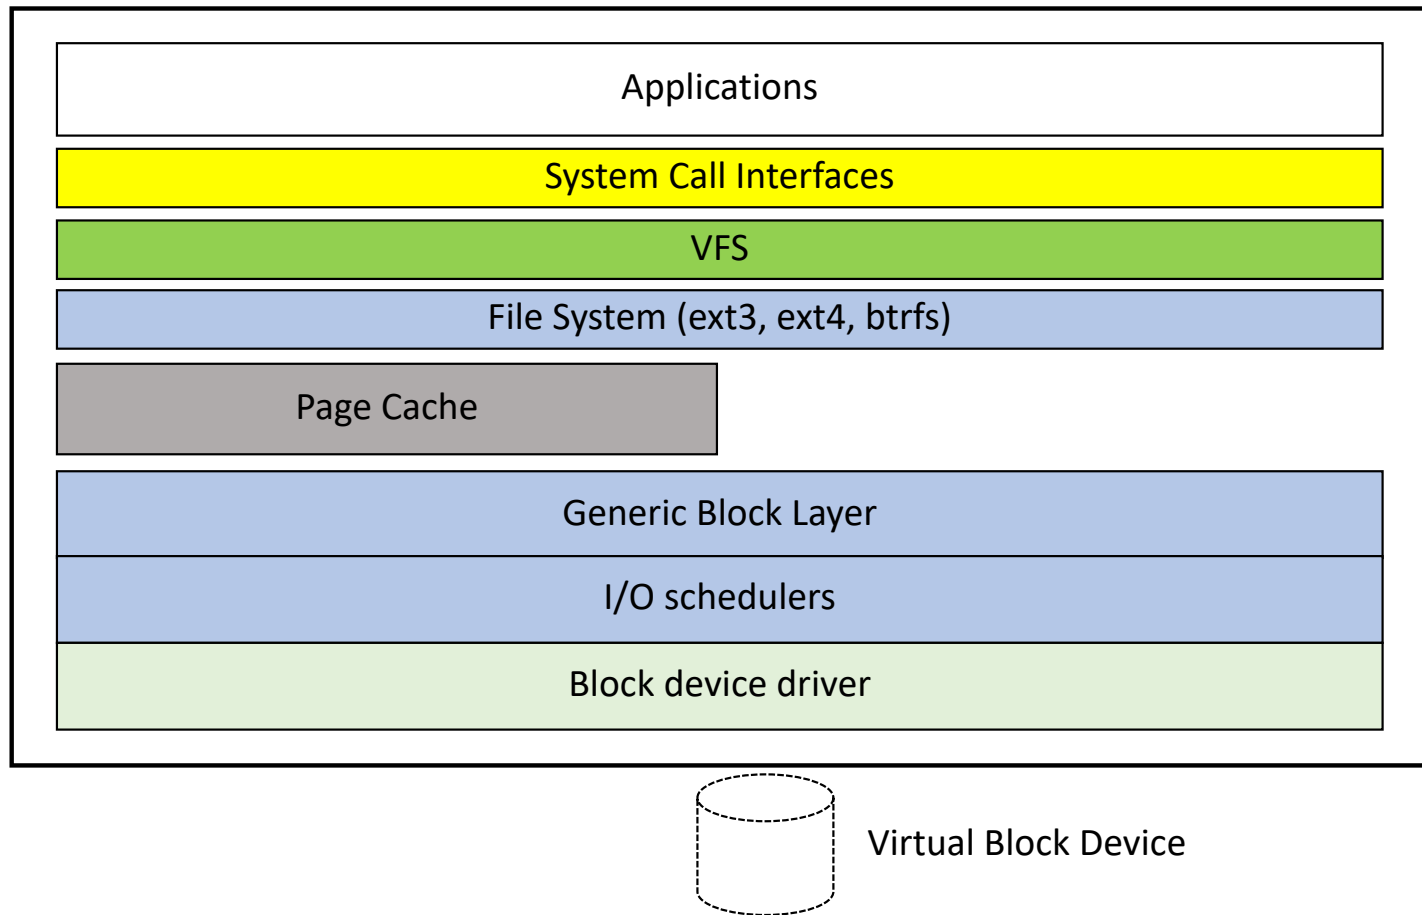


# I/O layers within a single host

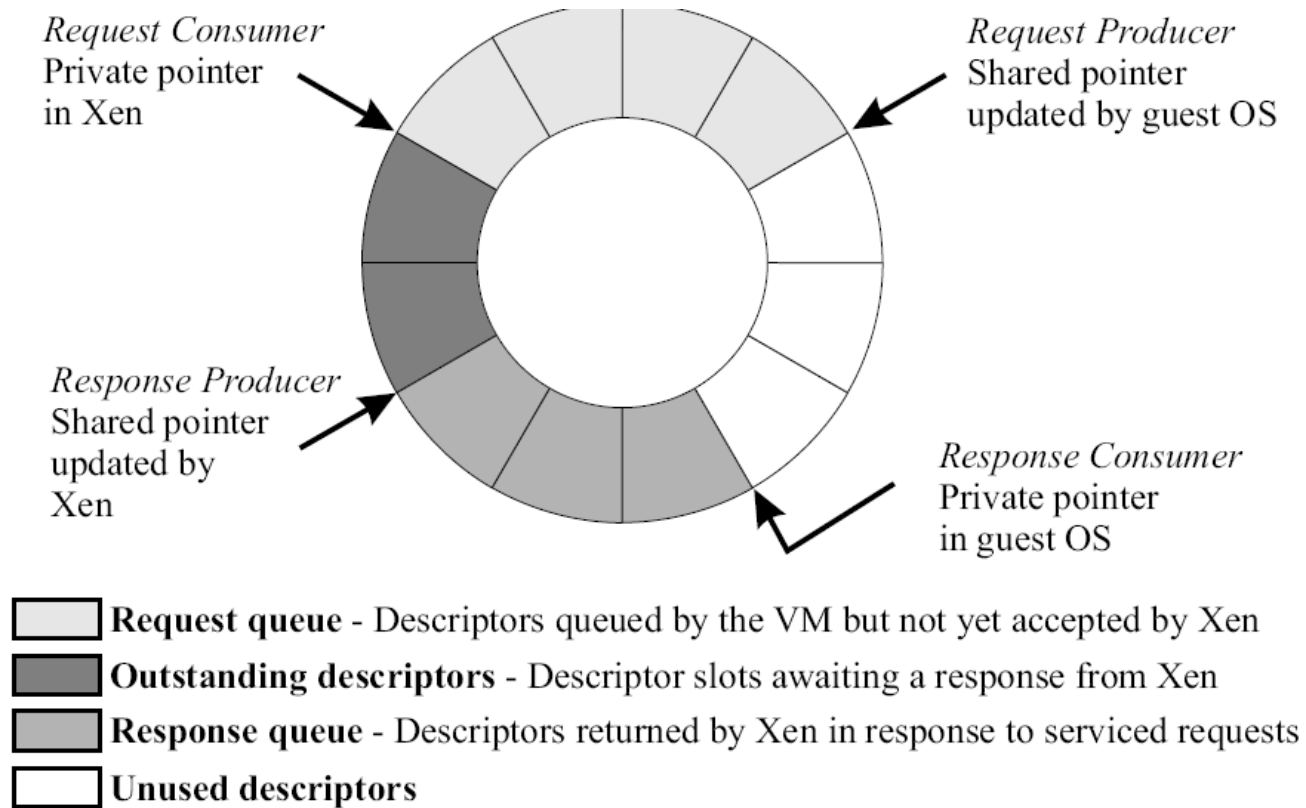


# I/O layers within a VM

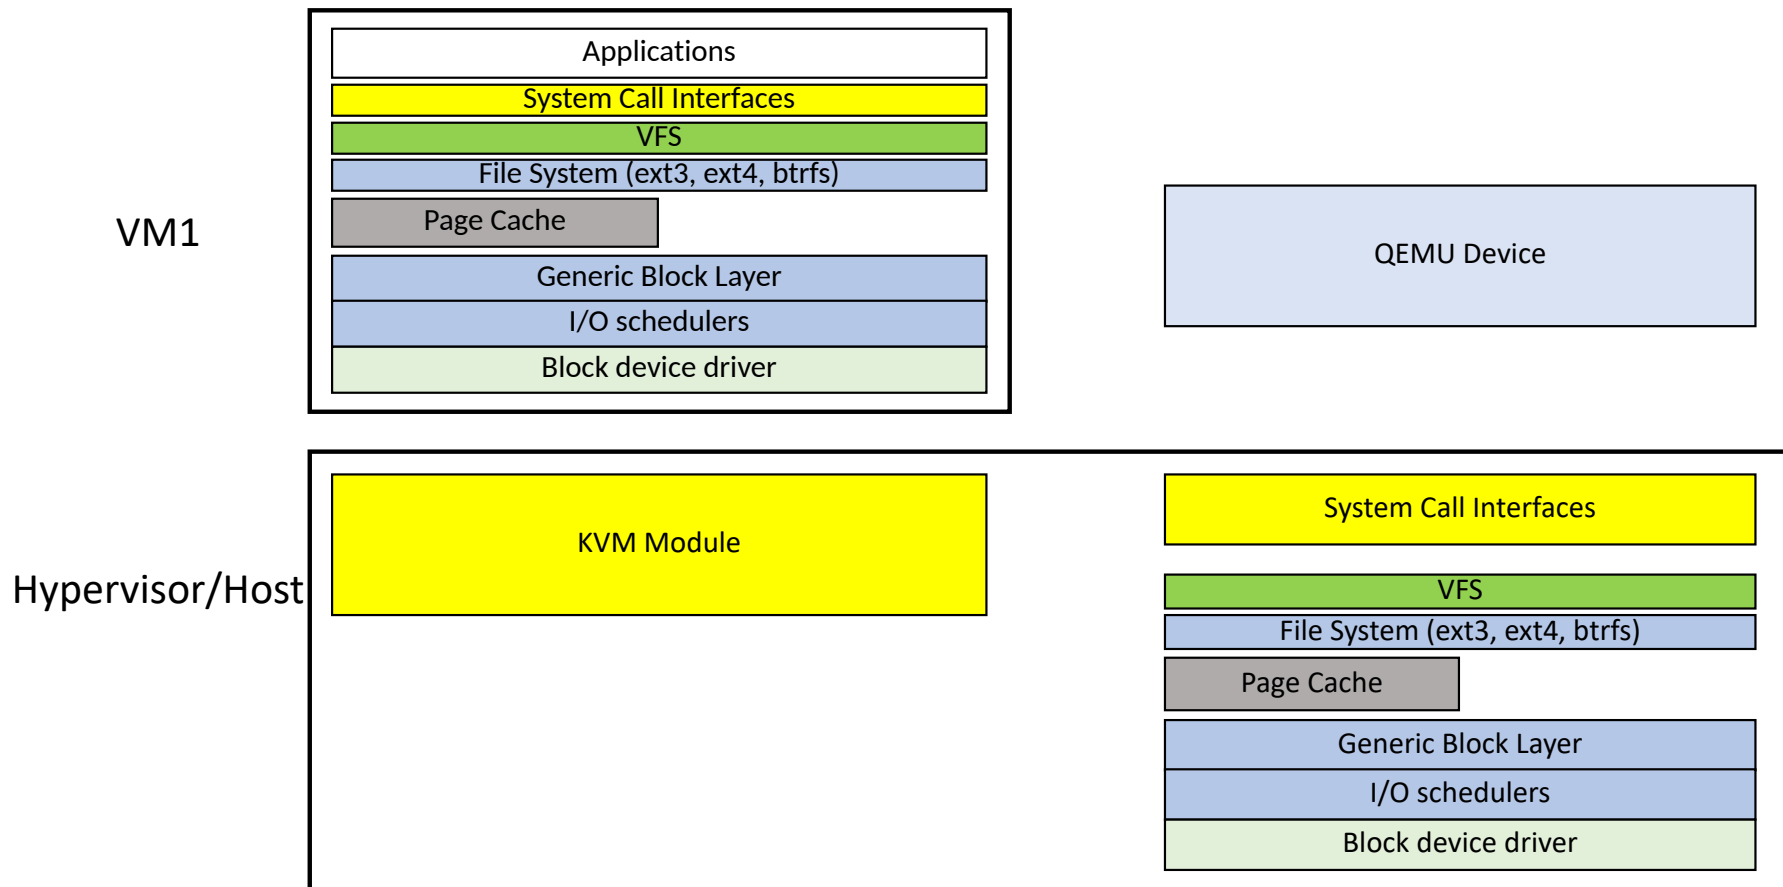
VM1



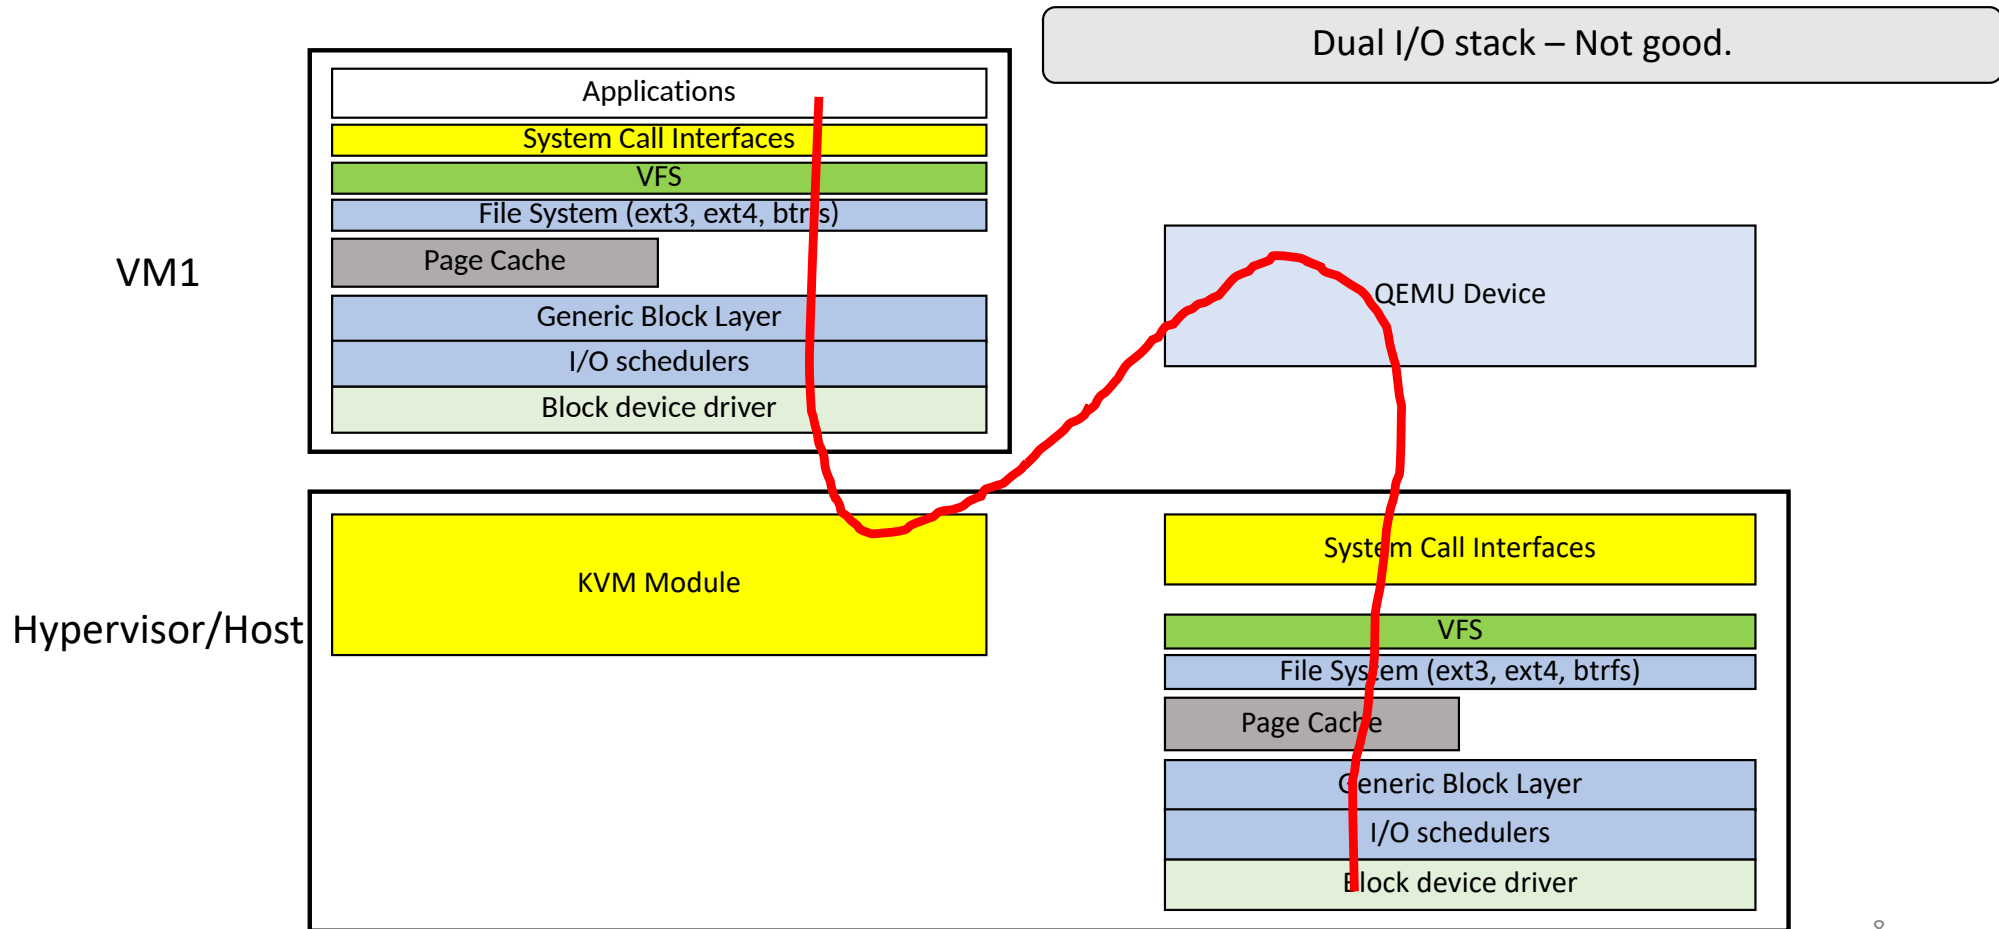
# I/O Rings



# I/O layers in Virtualization

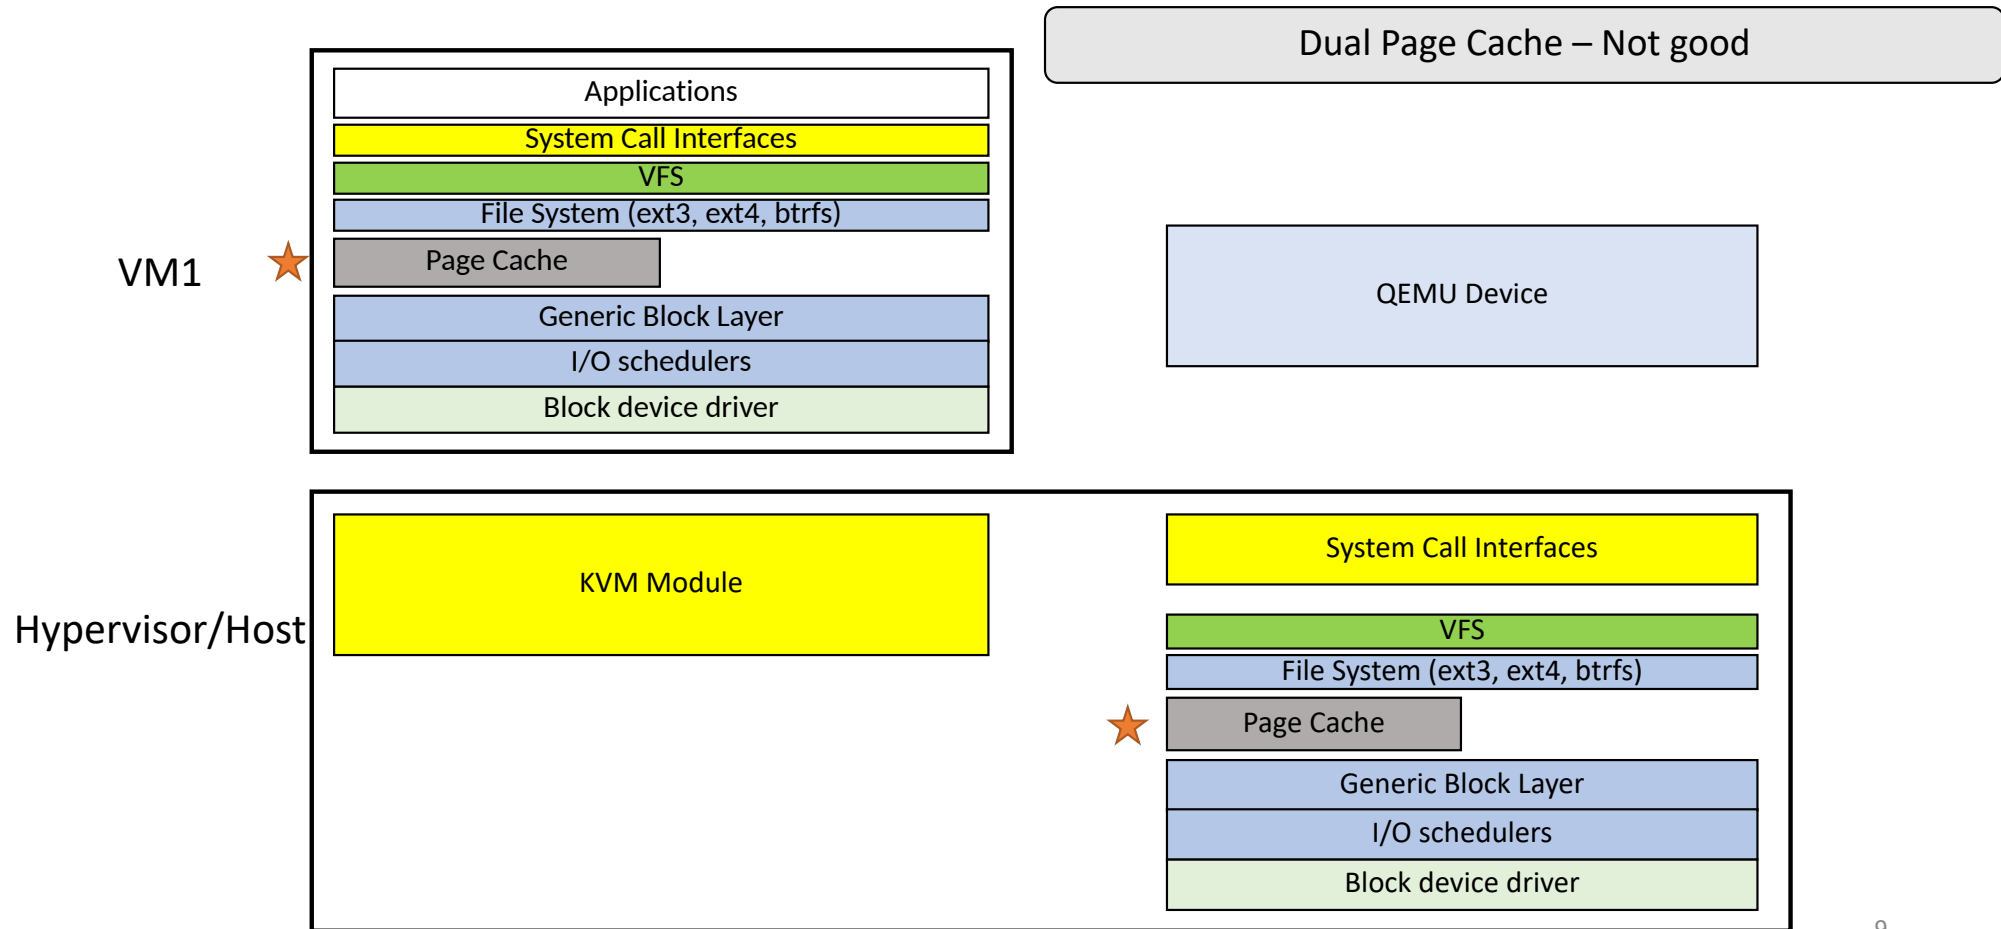


# I/O layers in Virtualization






# I/O layers in Virtualization




# I/O Data Plane Redundancy



Multiple data copying steps for data communication between two VMs.  
Not good!

# Virtualize Storage Device



- Virtual disk is stored as
  - a file in the host file system
  - or partition on physical disk
- Operations to the block device is emulated by QEMU
- Guest issues block reads & writes
- QEMU converts them to file operations on the virtual disk file

# Virtual Disk Image Type Matters!

- A “pre-allocated” disk image (1 virtual to 1 physical block)
  - A 10 GB disk image reserves 10 GB of disk space, regardless of whether the virtual machine guests uses 1 GB or 10 GB (allocated at creation time)
- An “extensible” disk image, useful for growing on demand
  - From the VM point of view, it sees a full size disk, but the hypervisor is actually lying to the VM, and is allocating the disk blocks on the HOST side on demand

# Disk images - pros / cons

- A “pre-allocated” disk image
  - Pros: Fast
  - Cons: Uses all space
- An extensible disk image
  - Pros: Less space
  - Cons: A bit overhead, fragmentation
- It depends on what we are trying to achieve: system design tradeoff

# VM Creation and Virtual Disk Images

- Assume that each virtual machine (VM) needs a disk image. If we are only going to create a single VM, it's easy:
  - Create VM
    - (1) create disk image
    - (2) attach ISO image (installation) to start VM
    - (3) install operating system
    - (4) Done!
- What if we want to install 2 VMs ? We could probably install a second time. What about when we have to build 5 ? 40 ? And do this very often (e.g., cloud service vendors)?
  - How do you increase the efficiency of such VM creation?

# Two Concrete Techniques

- Raw disks (“pre-allocated”)
  - Byte-for-byte disk image, byte 0 = byte 0 of the disk
- QEMU-KVM’s “QCOW2” (Qemu Copy On Write, v.2) format (extensible)
  - Grow-on-demand
  - Compression support
  - Encryption support
  - Copy-on-write!

# What is Copy-on-Write?


- Traditionally (e.g., raw disks):
  - When programs inside the guest VM write to the virtual disk, the changes are written to the disk image in place.
- Copy-on-write:
  - Write delta and store somewhere else (don't modify the original copy)



# Use of CoW

- A new disk image, originates from a “master” image as a backing file.
  - E.g, `qemu-img create -o backing_file=master_image.qcow2 -f guest1.qcow2 10G`
- Initially, the size of
  - `guest1.qcow2` is 0 bytes.
  - backing file (`master_image.qcow2`) is (say) 10 GiB.
- For writes, KVM will write the changes to the `guest1.qcow2`. The file `master_image.qcow2` is never written to.
- For reads, KVM will read the block from the `master_image.qcow2` or `guest1.qcow2` (whichever is latest).

# Cloud Block Storage System




# Storage Area Network (SAN)

- Dedicated high-speed network interconnects and presents shared pools of storage devices to servers. (e.g., Fibre Channel)
- Light-weight solution: Protocols: iSCSI – Reuse Ethernet Network (by encapsulating [SCSI](#) commands into IP packets that don't require an FC connection)



# iSCSI

- iSCSI is a Storage Area Network (SAN) protocol that allows for SCSI command transmission over a TCP/IP network
- iSCSI allows for the sharing of I/O devices over a long distance.
- iSCSI maintains the SCSI notion of an Initiator and Target device



# Data Deduplication

- Duplicate data is deleted leaving, only one copy of the data to be stored.
- Compare new data block to existing data blocks.
  - If contents of new block are unique then store it in the disk.
  - But if it is a duplicate of existing blocks then don't store again but create a reference.
- Only one unique instance of the data is retained on storage media (e.g., disk). Redundant data is replaced with a pointer to the unique data copy.

# Deduplication Methods

- In-line deduplication:
  - Hash calculations are created as the data is entered in real time.
  - If the target device identifies a block that has already been stored then it simply references to the existing block.
- Pros: Inline deduplication significantly reduces the raw disk capacity needed in the system since the full, not-yet-deduplicated data set is never written to disk
- Cons: However, “because hash calculations and lookups takes so long, data writes can be slower thereby reducing the backup throughput of the device.”
- What is off-line deduplication?